

FIG. 1

HALO TRACEABILITY HARNESS — FULL APPARATUS

100 Halo Traceability Harness




FIG. 2

TAGGER SUB-APPARATUS (expansion of element 108)

Claim Candidate A — invariant per-claim four-surface schema





FIG. 3

SCORING MODULE — PRINCIPLE-INDEXED REDUCTION (expansion of 112)

Claim Candidate B — framework-indexed status matrix




Distinguishing feature vs. existing LLM-eval frameworks: output is a framework-indexed status matrix, not a single benchmark score; structurally a different artifact.

FIG. 4

METHOD / DATA-FLOW SEQUENCE — steps S1 through S8

Claim Candidate C — evaluator-side end-to-end method; Deployment never cooperates



All steps S1-S8 are performed by the apparatus 100. Deployment 106 lies outside 101 and contributes only verbatim public-API responses. No SDK install, log access, weight access, or instrumentation hook on 106 is required at any step. Audit Store 110 retains every tagged_interaction_record for independent replay against the same model checkpoint.